
Distributed Memory Programming Model

Course: Parallel and Distributed Computing

Prepared by: Maryam Saif



Here's a clean, simple explanation of the **Distributed Memory Programming Model** that will help you understand it clearly (especially for parallel & high-performance computing):

Distributed Memory Programming Model — Simple Explanation

1. What It Is

In a **distributed memory system**, each processor (or node) has **its own private memory**. No two processors share memory directly.

To exchange data, they must **explicitly send and receive messages**.

Think of it like multiple computers connected via a network, each with its own RAM.

2. How It Works

Each processor runs its own copy of the program and:

- Stores its own local variables/data
- Performs computations independently
- Communicates with other processors **only when needed** through message passing

Communication happens using libraries like:

- **MPI (Message Passing Interface)** → Most common
 - PVM (Parallel Virtual Machine)
-

3. Key Operations

In message passing (MPI-style), you commonly use:

✓ Send

A processor sends data to another processor.

✓ Receive

A processor receives data from another processor.

✓ Collective Communication

These operations involve **all processors**:

- Broadcast (one → many)
 - Scatter (one → parts to many)
 - Gather (many → one)
 - All-to-all (everyone sends to everyone)
 - Reduction (e.g., sum, min, max from all nodes)
-

4. Characteristics

Feature	Distributed Memory
Memory sharing	No shared memory

Feature	Distributed Memory
Communication	Explicit message passing
Data access	Local only, need send/receive for remote
Scalability	Very high (thousands of nodes)
Typical systems	Clusters, supercomputers

5. Advantages

- ✓ Very scalable
 - ✓ No memory contention between processors
 - ✓ Works well for very large data sets
-

6. Disadvantages

- ✗ Harder to program
 - ✗ Need to manually handle communication
 - ✗ Debugging message passing is difficult
-

7. Simple Example

Suppose you want processors to share their computed value:

Processor 0 wants to send value x to Processor 1:

```
P0: MPI_Send(x, destination=1)
P1: MPI_Recv(x, source=0)
```

Everything must be done explicitly.

8. Where It Is Used

- Scientific simulations
- Weather prediction
- Machine learning clusters
- Large matrix computations

- HPC (High Performance Computing)
-
-

1. Introduction

- A memory model defines how processors access memory.
- Shared memory: all processors share same memory.
- Distributed memory: each processor has private memory, data sharing via message passing.

Here's a polished, clear **introduction paragraph** based on what you wrote — clean, academic, and suitable for an assignment or exam:

Introduction

A **memory model** describes the way processors access and manage memory during parallel execution. In the **shared memory model**, all processors operate on a common global memory space, allowing them to directly read and write shared data. In contrast, the **distributed memory model** assigns each processor its own private memory, requiring data exchange through explicit **message passing**. These models shape how parallel programs are designed, synchronized, and optimized.

2. Concept of Distributed Memory

- Each processor/node has its own local memory.
- No direct sharing of data — communication via messages.
- Local data: stored in node's own memory.
- Remote data: accessed via network messages.
- Analogy: Chefs in separate kitchens send messages to borrow ingredients!

Here's a clean, well-structured and academically polished version of your content:

2. Concept of Distributed Memory

In a distributed memory system, each processor or node maintains its **own local memory**, which cannot be directly accessed by other processors. This means there is **no global shared memory space**. Any exchange of data between processors occurs through **explicit message passing** over a network.

Local data refers to the information stored within a processor's private memory, while remote data must be accessed by sending and receiving messages across connected nodes. A useful analogy is imagining multiple chefs working in **separate kitchens**; since they cannot access each other's ingredients directly, they must **send messages** to request or share items.

3. Architecture Overview

- Independent nodes each with CPU + Memory + Storage.
- Connected by high-speed network (Ethernet, InfiniBand).
- Communication handled via explicit message passing.
- Example: HPC clusters, cloud systems.

3. Architecture Overview

Distributed memory systems consist of multiple **independent nodes**, each equipped with its own CPU, memory, and storage. These nodes do not share physical memory and operate autonomously. They are interconnected through **high-speed networks** such as Ethernet or InfiniBand, enabling communication between them.

Since memory is not shared, all inter-node coordination relies on **explicit message passing**, where processors send and receive data using communication protocols. This architectural design is commonly seen in **high-performance computing (HPC) clusters**, large-scale servers, and modern cloud computing systems.

4. Communication in Distributed Memory

- Processes communicate using message passing libraries:
 - MPI (Message Passing Interface)
 - PVM (Parallel Virtual Machine)
 - Sockets / RPC
 - Communication types:
 - Point-to-point
 - Collective (broadcast, scatter, gather)
-

4. Communication in Distributed Memory

In distributed memory systems, processes exchange data through **explicit message passing**, since no shared memory exists between nodes. This communication is implemented using specialized libraries and protocols such as:

- **MPI (Message Passing Interface)** – the most widely used standard in HPC
- **PVM (Parallel Virtual Machine)** – an older but still relevant system for distributed computing
- **Sockets / RPC (Remote Procedure Calls)** – low-level mechanisms for custom communication

Communication typically occurs in two main forms:

- **Point-to-Point Communication:** Direct data exchange between two specific processes using send/receive operations.
 - **Collective Communication:** Involves a group of processes participating in coordinated communication patterns such as **broadcast**, **scatter**, **gather**, **all-to-all**, and **reduction**.
-

5. Example: MPI Communication

- Example code:
 - if (rank == 0): send(data, 1)
 - else if (rank == 1): recv(data, 0)
- • MPI_Send / MPI_Recv for data exchange.
- • Blocking or non-blocking communication possible.

5. Example: MPI Communication

A basic illustration of message passing in MPI involves two processes exchanging data. Each process is identified by its **rank**, and communication is performed using explicit send and receive operations. For example:

```
if (rank == 0) {  
    MPI_Send(data, destination=1);  
}  
else if (rank == 1) {  
    MPI_Recv(data, source=0);  
}
```

MPI provides several communication routines, with **MPI_Send** and **MPI_Recv** being the most fundamental for point-to-point data exchange. These operations can be either **blocking**, where a process waits until the operation completes, or **non-blocking**, allowing a process to continue executing while communication occurs in the background.

6. Data Distribution Strategies

- Block distribution – equal blocks per process.
- Cyclic distribution – every n th element.
- Block-cyclic – hybrid for better load balance.
- Example: Matrix multiplication across 4 processors.

6. Data Distribution Strategies

Efficient data distribution is essential in distributed memory systems to ensure balanced workloads and minimize communication overhead. Common strategies include:

- **Block Distribution:**
Data is divided into large, contiguous blocks, and each process receives one block. This approach is simple and works well when all processes perform similar amounts of work.
- **Cyclic Distribution:**
Elements are assigned to processes in a round-robin fashion (e.g., every n th element). This ensures better balancing when workloads vary across data elements.
- **Block-Cyclic Distribution:**
A hybrid strategy where small blocks of data are distributed cyclically. It provides both locality (from blocks) and improved load balancing (from cyclic distribution).

A typical use case is **parallel matrix multiplication**, where a matrix is partitioned across four processors using block or block-cyclic schemes to distribute computation evenly.

7. Programming Models and Tools

- MPI – Standard library for distributed memory systems.
 - PVM – Older, used for heterogeneous environments.
 - MapReduce, Spark – High-level distributed data frameworks.
 - Cloud frameworks (AWS, Kubernetes) – use distributed memory underneath.
-

7. Programming Models and Tools

A variety of programming models and tools support the development of applications for distributed memory systems:

- **MPI (Message Passing Interface):**
The industry-standard library for explicit message passing, widely used in scientific computing and HPC environments.
 - **PVM (Parallel Virtual Machine):**
An older system designed for coordinating heterogeneous networked computers, still relevant in some legacy applications.
 - **MapReduce and Apache Spark:**
High-level distributed data processing frameworks that abstract away low-level message passing and provide simplified models for large-scale data analytics.
 - **Cloud-Based Frameworks (AWS, Kubernetes, etc.):**
Modern cloud infrastructures rely on distributed memory principles, using clusters of independent nodes to run scalable containerized or serverless applications.
-

8. Advantages and Disadvantages

- Advantages:
 - • Highly scalable
 - • No memory contention
 - • Fault isolation

- Disadvantages:
 - • Complex programming and debugging
 - • Communication overhead
 - • Manual data partitioning required

8. Advantages and Disadvantages

Advantages:

- **Highly scalable:** Additional nodes can be added easily, making the system suitable for large-scale parallel applications.
- **No memory contention:** Each processor has its own local memory, eliminating conflicts from simultaneous access.
- **Fault isolation:** Failure of one node does not directly affect the memory or execution of other nodes.

Disadvantages:

- **Complex programming and debugging:** Developers must explicitly manage communication, synchronization, and data distribution.
 - **Communication overhead:** Data exchange across the network can introduce latency and impact performance.
 - **Manual data partitioning:** Programmers must divide data and tasks among processors, increasing development effort.
-

9. Real-Life Examples

- • Scientific simulations (weather forecasting)
 - • Big data analytics (Hadoop, Spark)
 - • Cloud computing (VMs with separate memory)
 - • Online gaming servers (regional data exchange)
-
- Analogy: Game regions communicating player data.

9. Real-Life Examples

Distributed memory systems are widely used in applications that require parallel processing and large-scale data management:

- **Scientific Simulations:** Weather forecasting, climate modeling, and computational physics rely on multiple nodes working together on different parts of the problem.
- **Big Data Analytics:** Frameworks like Hadoop and Apache Spark distribute data across nodes to process massive datasets efficiently.
- **Cloud Computing:** Virtual machines and containers operate with separate memory spaces while communicating over the network.
- **Online Gaming Servers:** Regional servers exchange player data to maintain consistent game states across locations.

Analogy: Similar to **different regions in an online game** communicating player information, each node handles its own local tasks while sharing necessary data with others.

10. Comparison: Shared vs Distributed Memory

- | Feature | Shared | Distributed |
 - |-----|-----|-----|
 - | Memory Space | Common | Private |
 - | Communication | Implicit | Explicit |
 - | Scalability | Low | High |
 - | Example | Multicore CPU | Cluster (MPI) |
-

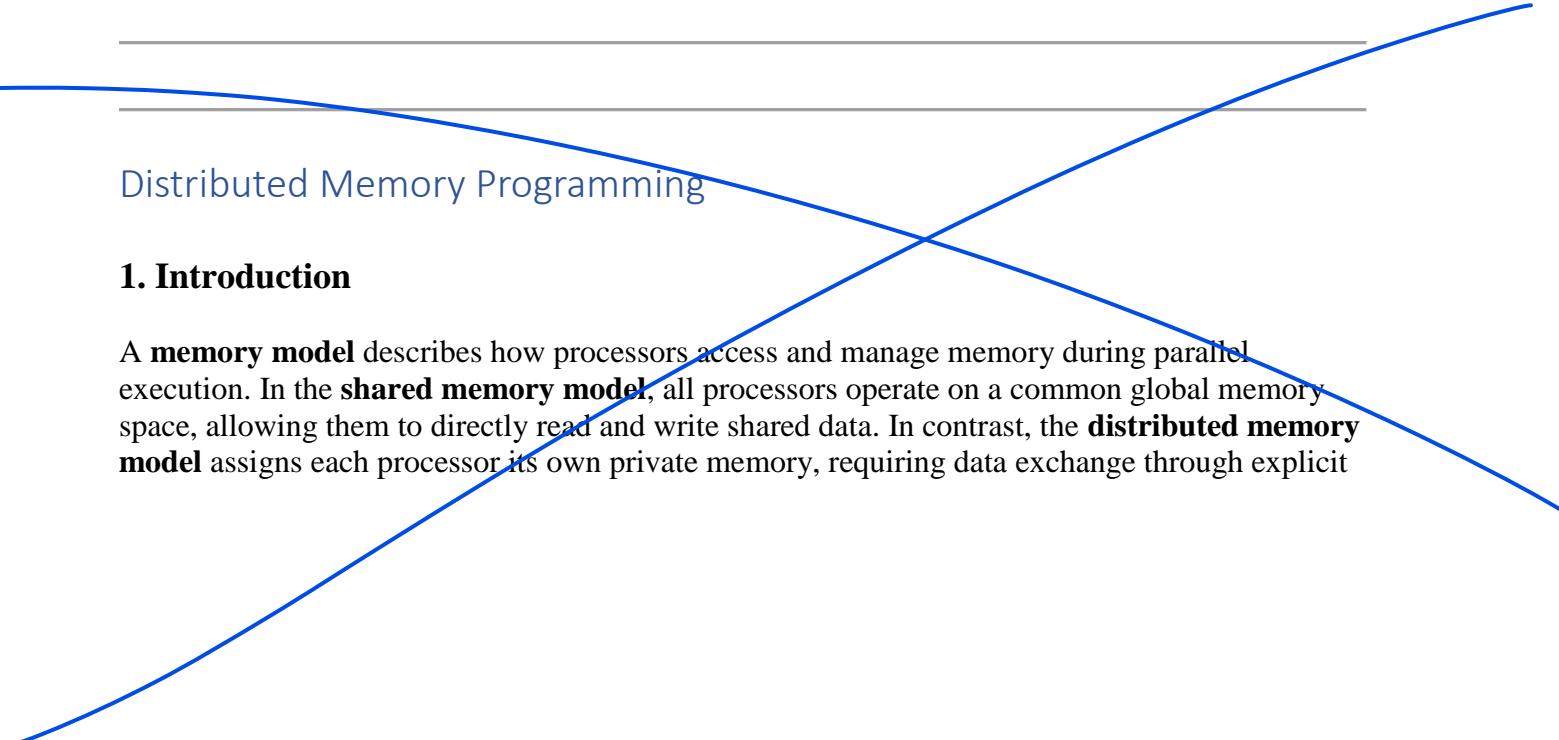
10. Comparison: Shared vs Distributed Memory

Feature	Shared Memory	Distributed Memory
Memory Space	Common, accessible by all processors	Private to each processor/node
Communication	Implicit (via shared memory)	Explicit (via message passing)
Scalability	Limited (contention issues)	High (nodes can be added easily)
Example	Multicore CPU systems	Clusters using MPI or cloud nodes

11. Summary

- • Distributed memory = independent nodes + explicit communication.
 - • Core library: MPI.
 - • Used in HPC, clusters, and cloud computing.
 - • High scalability with communication challenges.
-

11. Summary

- **Distributed memory systems** consist of independent nodes, each with its own local memory, and require **explicit communication** for data exchange.
 - The **Message Passing Interface (MPI)** is the core library widely used for programming these systems.
 - Distributed memory architectures are common in **high-performance computing (HPC)**, **clusters**, and **cloud computing** environments.
 - They offer **high scalability**, but programming complexity and communication overhead remain key challenges.
-
- 

Distributed Memory Programming

1. Introduction

A **memory model** describes how processors access and manage memory during parallel execution. In the **shared memory model**, all processors operate on a common global memory space, allowing them to directly read and write shared data. In contrast, the **distributed memory model** assigns each processor its own private memory, requiring data exchange through explicit

message passing. These models shape how parallel programs are designed, synchronized, and optimized.

2. Concept of Distributed Memory

In a distributed memory system, each processor or node maintains its **own local memory**, which cannot be directly accessed by other processors. This means there is **no global shared memory space**. Any exchange of data between processors occurs through **explicit message passing** over a network.

Local data refers to the information stored within a processor's private memory, while remote data must be accessed by sending and receiving messages across connected nodes.

Analogy: Multiple chefs working in **separate kitchens** must **send messages** to request or share ingredients since they cannot access each other's supplies directly.

3. Architecture Overview

Distributed memory systems consist of multiple **independent nodes**, each equipped with its own CPU, memory, and storage. These nodes are interconnected through **high-speed networks** such as Ethernet or InfiniBand. Since memory is not shared, all inter-node coordination relies on **explicit message passing**, enabling scalable computation.

Examples: High-performance computing (HPC) clusters, large-scale servers, and modern cloud computing systems.

4. Communication in Distributed Memory

Processes exchange data through **explicit message passing** using specialized libraries:

- **MPI (Message Passing Interface)** – industry standard for HPC
- **PVM (Parallel Virtual Machine)** – older system for heterogeneous networks
- **Sockets / RPC (Remote Procedure Calls)** – low-level communication methods

Types of Communication:

- **Point-to-Point:** Direct send/receive between two processes.
- **Collective:** Coordinated operations among multiple processes (broadcast, scatter, gather, all-to-all, reduction).

5. Example: MPI Communication

A simple MPI example with two processes exchanging data:

```
if (rank == 0) {
    MPI_Send(data, destination=1);
}
else if (rank == 1) {
    MPI_Recv(data, source=0);
}
```

- **MPI_Send / MPI_Recv** are fundamental for point-to-point communication.
- Operations can be **blocking** (wait until complete) or **non-blocking** (continue execution while communication occurs).

6. Data Distribution Strategies

Efficient data distribution ensures balanced workloads and minimizes communication overhead:

- **Block Distribution:** Divide data into contiguous blocks; each process gets one block.
- **Cyclic Distribution:** Assign elements to processes in a round-robin fashion.
- **Block-Cyclic Distribution:** Hybrid approach combining blocks and cyclic assignment for better load balance.

Example: Parallel matrix multiplication across four processors using block or block-cyclic distribution.

7. Programming Models and Tools

Tools and frameworks for distributed memory systems include:

- **MPI:** Standard library for explicit message passing.
 - **PVM:** Legacy system for heterogeneous networks.
 - **MapReduce / Apache Spark:** High-level frameworks abstracting low-level message passing.
 - **Cloud Frameworks (AWS, Kubernetes):** Underlying distributed memory principles used for scalable services.
-

8. Advantages and Disadvantages

Advantages:

- **Highly scalable** – nodes can be added easily.
- **No memory contention** – each processor has private memory.
- **Fault isolation** – node failure does not directly affect others.

Disadvantages:

- **Complex programming and debugging** – explicit communication required.
 - **Communication overhead** – network latency can impact performance.
 - **Manual data partitioning** – programmer must divide tasks and data among nodes.
-

9. Real-Life Examples

- **Scientific Simulations:** Weather forecasting, climate modeling.
- **Big Data Analytics:** Hadoop and Spark distribute data for large-scale processing.
- **Cloud Computing:** Virtual machines and containers use separate memory spaces.
- **Online Gaming Servers:** Regional servers exchange player data to maintain consistent game states.

Analogy: Different regions in an online game communicate player information, with each region managing its own tasks locally.

10. Comparison: Shared vs Distributed Memory

Feature	Shared Memory	Distributed Memory
Memory Space	Common, accessible by all processors	Private to each processor/node
Communication	Implicit (via shared memory)	Explicit (via message passing)
Scalability	Limited (contention issues)	High (nodes can be added easily)
Example	Multicore CPU systems	Clusters using MPI or cloud nodes

11. Summary

- **Distributed memory systems** consist of independent nodes, each with local memory, requiring **explicit communication**.
 - **MPI** is the core library for programming these systems.
 - Commonly used in **HPC, clusters, and cloud computing**.
 - Provide **high scalability**, but involve programming complexity and communication overhead.
-