

CUDA Programming Introduction

Andreas Moshovos

Winter 2009

Some slides/material from:

UIUC course by Wen-Mei Hwu and David Kirk


UCSB course by Andrea Di Blas

Universitat Jena by Waqar Saleem

NVIDIA by Simon Green

Programmer's view

- GPU as a co-processor



Target Applications

```
int a[N]; // N is large  
for all elements of a compute  
    a[i] = a[i] * fade
```

- Lots of **independent** computations
 - CUDA thread need not be independent

Programmer's View of the GPU

- GPU: a compute **device** that:
 - Is a coprocessor to the CPU or **host**
 - Has its own DRAM (**device memory**)
 - Runs many **threads** in parallel
- Data-parallel portions of an application are executed on the device as **kernels** which run in parallel on many threads

Why are threads useful?


- Concurrency:
 - Do multiple things in parallel
- 
- Needs more functional units

- Put more hardware → Get higher performance


GPU vs. CPU Threads

- GPU threads are extremely lightweight
 - Very little creation overhead
 - In the order of microseconds
- GPU needs 1000s of threads for full efficiency
 - Multi-core CPU needs only a few

Execution Timeline



GBT: Grids of Blocks of Threads




Why? Realities of integrated circuits: need to cluster computation and storage to achieve high speeds

Grids of Blocks of Threads: Dimension Limits

- Grid of Blocks 1D or 2D
 - Max x: 65535
 - Max y: 65535
- Block of Threads: 1D, 2D, or 3D
 - Max number of threads: 512
 - Max x: 512
 - Max y: 512
 - Max z: 64
- Limits apply to Compute Capability 1.0, 1.1, 1.2, and 1.3
 - GTX280 = 1.3

Block and Thread IDs


- Threads and blocks have IDs
 - So each thread can decide what data to work on
 - Block ID: 1D or 2D
 - Thread ID: 1D, 2D, or 3D
- Simplifies memory addressing when processing multidimensional data




- IDs and dimensions are easily accessible through predefined “variables”, e.g., `blockDim.x` and `threadIdx.x`

Thread Batching

- A kernel is executed as a grid of thread blocks
 - All threads share data memory space
- A **thread block**: threads that can **cooperate** with each other by:
 - Synchronizing their execution
 - For hazard-free shared memory accesses
 - Efficiently sharing data through a low latency **shared memory**
- Two threads from two different blocks cannot cooperate



Programmer's view: Memory Model



Programmer's View: Memory Detail – Thread and Host

- Each thread can:
 - R/W per-thread registers
 - R/W per-thread local memory
 - R/W per-block shared memory
 - R/W per-grid global memory
 - Read only per-grid constant memory
 - Read only per-grid texture memory
- The host can R/W:
 - global, constant, and texture memories

Memory Model: Global, Constant, and Texture Memories

- Global memory
 - Main means of communicating R/W Data between host and device
 - Contents visible to all threads
- Texture and Constant Memories
 - Constants initialized by host
 - Contents visible to all threads

Memory Model Summary


Memory	Location	Cached	Access	Scope
Local	off-chip	No	R/W	thread
Shared	on-chip	N/A	R/W	all threads in a block
Global	off-chip	No	R/W	all threads + host
Constant	off-chip	Yes	RO	all threads + host
Texture	off-chip	Yes	RO	all threads + host

- **Execution order is undefined**
- Do not assume and use:
 - block 0 executes before block 1
 - Thread 10 executes before thread 20
 - And **any** other ordering even if you can observe it

Execution Model Summary

- **Grid of blocks of threads**
 - 1D/2D grid of blocks
 - 1D/2D/3D blocks of threads
- **All blocks are identical:**
 - same structure and # of threads
- **Block execution order is undefined**
- Same block threads:
 - can synchronize and share data fast (shared memory)
- Threads from different blocks:
 - Cannot cooperate
 - Communication through global memory
- **Threads and Blocks have IDs**
 - Simplifies data indexing
 - Can be 1D, 2D, or 3D (threads)
- Blocks do not migrate: execute on the same processor
- Several blocks may run over the same processor

CUDA Software Architecture



Reasoning about CUDA call ordering

- GPU communication via `cuda...()` calls and kernel invocations
 - `cudaMalloc`, `cudaMemcpy`,
- Asynchronous from the CPU's perspective
 - CPU places a request in a “CUDA” queue
 - requests are handled in-order
- Streams allow for multiple queues
 - More on this much later one

```
int a[N];  
    for (i =0; i < N; i++)  
        a[i] = a[i] + x;
```

1. Allocate CPU Data Structure
2. Initialize Data on CPU
3. Allocate GPU Data Structure
4. Copy Data from CPU to GPU
5. Define *Execution Configuration*
6. Run Kernel
7. CPU synchronizes with GPU
8. Copy Data from GPU to CPU
9. De-allocate GPU and CPU memory

1. Allocate CPU Data

```
float *ha;

main (int argc, char *argv[])
{
    int N = atoi (argv[1]);
    ha = (float *) malloc (sizeof (float) * N);
    ...
}
```

- Pinned memory allocation results in faster CPU to/from GPU copies
- More on this later
- `cudaMallocHost (...)`

2. Initialize CPU Data

```
float *ha;  
  
int i;  
  
for (i = 0; i < N; i++)  
    ha[i] = i;
```

3. Allocate GPU Data

```
float *da;  
  
cudaMalloc ((void **) &da, sizeof (float) * N);
```

- Notice: no assignment side
 - NOT: da = cudaMalloc (...)
- Assignment is done internally:
 - That why we pass &da
- Allocated space in **Global Memory**

GPU Memory Allocation

- The host manages GPU memory allocation:
 - **cudaMalloc (void **ptr, size_t nbytes)**
 - Must explicitly cast to (**void ****)
 - `cudaMalloc ((void **) &da, sizeof (float) * N);`
 - **cudaFree (void *ptr) ;**
 - `cudaFree (da);`
 - **cudaMemset (void *ptr, int value, size_t nbytes);**
 - `cudaMemset (da, 0, N * sizeof (int));`
- Check the **CUDA Reference Manual**

4. Copy Initialized CPU data to GPU

```
float *da;  
float *ha;  
  
cudaMemcpy ((void *) da, // DESTINATION  
              (void *) ha, // SOURCE  
              sizeof (float) * N, // #bytes  
              cudaMemcpyHostToDevice); // DIRECTION
```

Host/Device Data Transfers

- The host initiates all transfers:
- `cudaMemcpy(void *dst, void *src,
size_t nbytes,
enum cudaMemcpyKind direction)`
- Asynchronous from the CPU's perspective
 - CPU thread continues
- In-order processing with other CUDA requests
- `enum cudaMemcpyKind`
 - `cudaMemcpyHostToDevice`
 - `cudaMemcpyDeviceToHost`
 - `cudaMemcpyDeviceToDevice`

5. Define Execution Configuration

- How many blocks and threads/block

```
int threads_block = 64;  
int blocks = N / threads_block;  
if (blocks % N != 0) blocks += 1;
```

- Alternatively:

```
blocks = (N + threads_block - 1) /  
         threads_block;
```

6. Launch Kernel & 7. CPU/GPU Synchronization

- Instructs the GPU to launch `blocks` x `thread_blocks` threads:

```
darradd <<<blocks, threads_block>> (da, 10f, N);  
cudaThreadSynchronize ();
```

- `darradd`: kernel name
- `<<<...>>>` execution configuration
 - More on this soon
- `(da, x, N)`: arguments
 - 256 – 8 byte limit / No variable arguments

CPU/GPU Synchronization

- CPU does not block on `cuda...()` calls
 - Kernel/requests are queued and processed in-order
 - Control returns to CPU immediately
- Good if there is other work to be done
 - e.g., preparing for the next kernel invocation
- Eventually, CPU must know when GPU is done
- Then it can safely copy the GPU results
- `cudaThreadSynchronize ()`
 - Block CPU until **all** preceding `cuda...()` and kernel requests have completed

8. Copy data from GPU to CPU & 9. DeAllocate Memory

```
float *da;  
float *ha;  
  
cudaMemcpy ((void *) ha, // DESTINATION  
              (void *) da, // SOURCE  
              sizeof (float) * N, // #bytes  
              cudaMemcpyDeviceToHost); // DIRECTION  
  
cudaFree (da);  
// display or process results here  
free (ha);
```

The GPU Kernel


```
__global__ darradd (float *da, float x, int N)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;

    if (i < N) da[i] = da[i] + x;
}
```

- **BlockIdx:** Unique Block ID.
 - Numerically ascending: 0, 1, ...
- **ThreadIdx:** Unique per Block Index
 - 0, 1, ...
 - Per Block
- **BlockDim:** Dimensions of Block
 - BlockDim.x, BlockDim.y, BlockDim.z
 - Unused dimensions default to 0

Array Index Calculation Example

```
int i = blockIdx.x * blockDim.x + threadIdx.x;
```



Assuming `blockDim.x = 64`