

# **CS-251: Parallel and Distributed Computing**

## **Lecture 04 – Hardware Architectures Shared Memory**

# Dr. Zeeshan Rafi

PhD MIS, MPhil IT, BS Software Engineering  
Former Software Engineer, Database Administrator  
System Analyst, Big Data Analyst  
Member Turkish Scientific & Technological Research Council

## Department of Computing and Information Technology

Istanbul University, TR



KHAS University, TR



University of Gujrat, PK



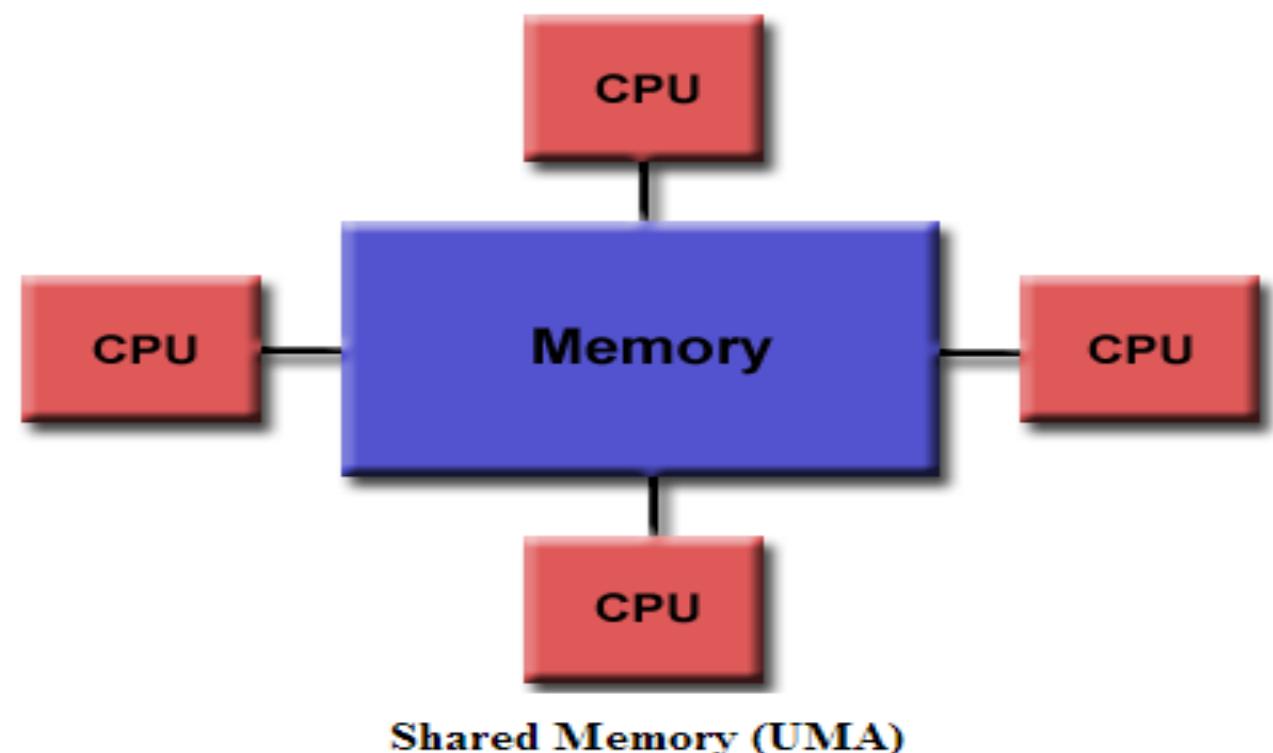
GCF University, Pk



# Shared Memory

# Shared Memory

- All processors access all memory as global address space
  - Multiple processors can operate independently but share the same memory resources
  - Changes in a memory location effected by one processor are visible to all other processors



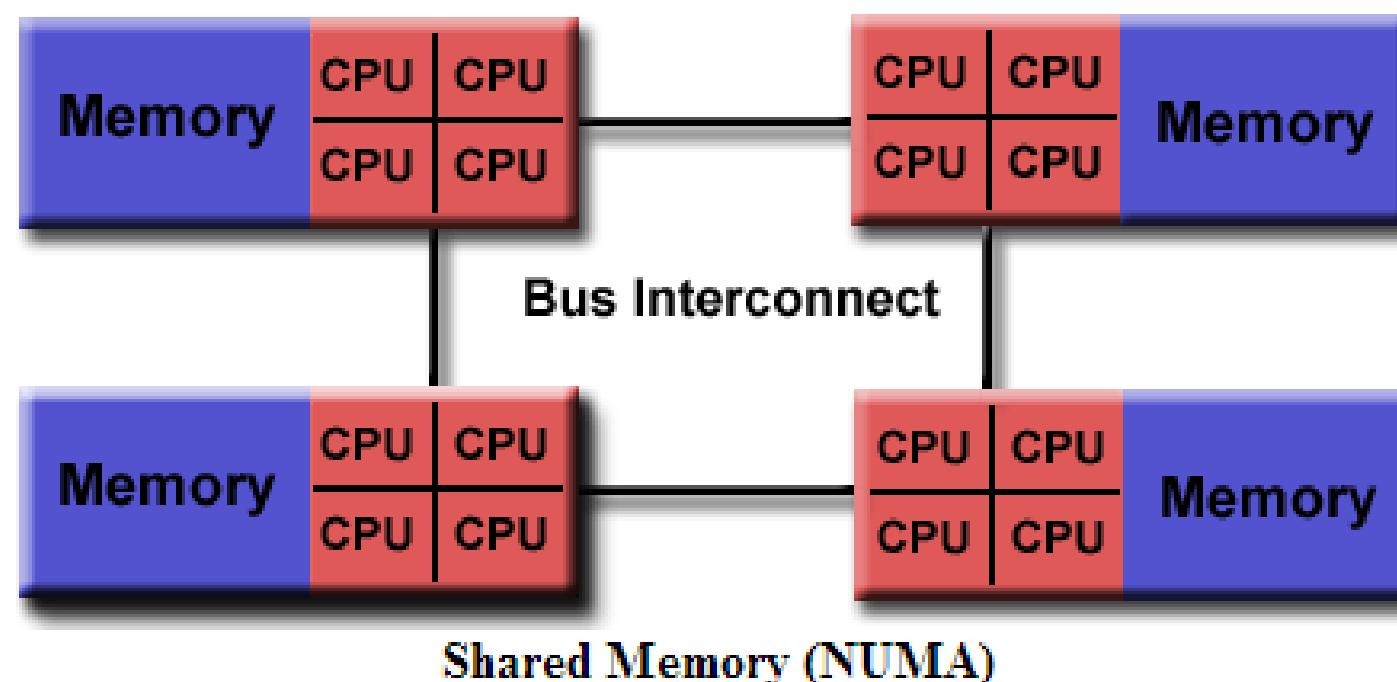
# Shared Memory

- Shared memory machines are **classified** as *UMA* and *NUMA*, based upon memory access times
  - Uniform Memory Access (UMA)
    - Most commonly represented today by Symmetric Multiprocessor (SMP) machines
    - Identical processors
    - Equal access and access times to memory
    - Sometimes called CC-UMA - Cache Coherent UMA
      - Cache coherent means if one processor updates a location in shared memory, all the other processors know about the update. Cache coherency is accomplished at the hardware level

# Shared Memory

## – Non-Uniform Memory Access (NUMA)

- Often made by physically linking two or more SMPs
- One SMP can directly access memory of another SMP
- Not all processors have equal access time to all memories
- Memory access across link is slower
- If cache coherency is maintained, then may also be called CC-NUMA - Cache Coherent NUMA

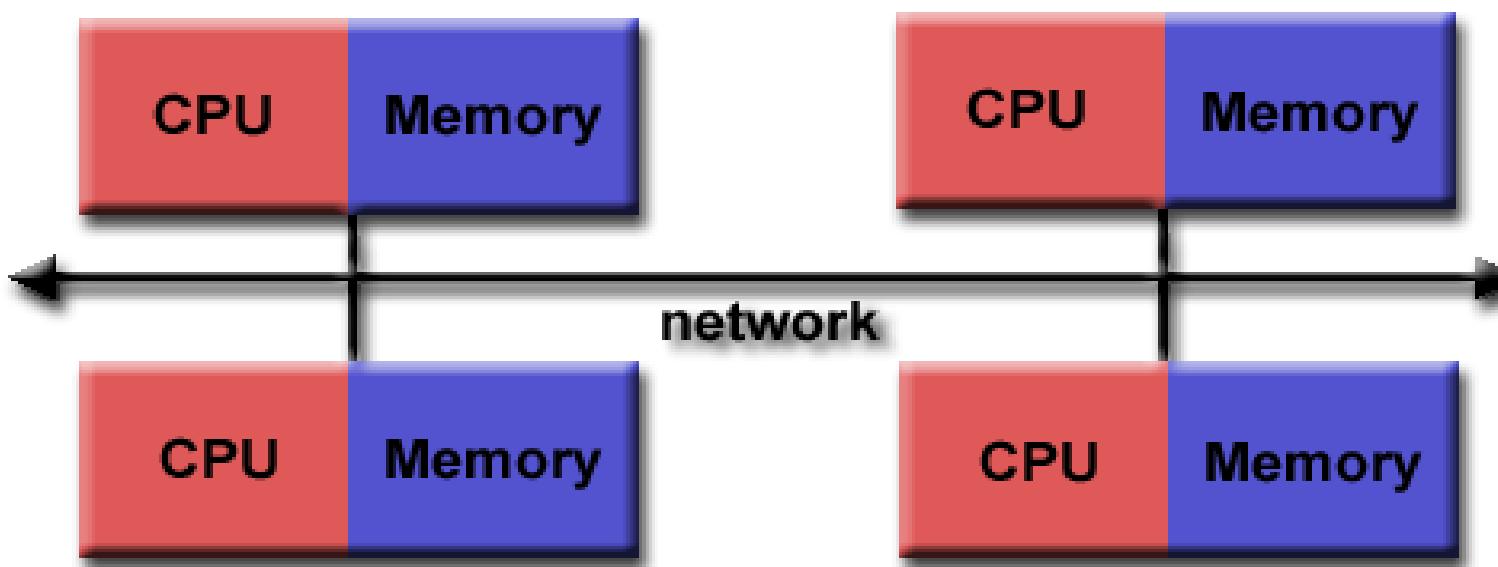


# Shared Memory

- Advantages
  - Global address space provides a user-friendly programming perspective to memory
  - Data sharing between tasks is both fast and uniform due to the proximity of memory to CPUs
- Disadvantages
  - Primary disadvantage is the lack of scalability between memory and CPUs
    - Adding more CPUs can geometrically increases traffic on the shared memory-CPU path, and for cache coherent systems, geometrically increase traffic associated with cache/memory management
  - Programmer responsibility for synchronization constructs that ensure "correct" access of global memory

# Distributed Memory

- Processors have their own **local memory**
  - Changes to processor's local memory have no effect on the memory of other processors
  - When a processor needs access to data in another processor, it is usually the task of the programmer to **explicitly define how and when data is communicated**
  - Synchronization between tasks is likewise the programmer's responsibility
  - The network "fabric" used for data transfer varies widely, though it

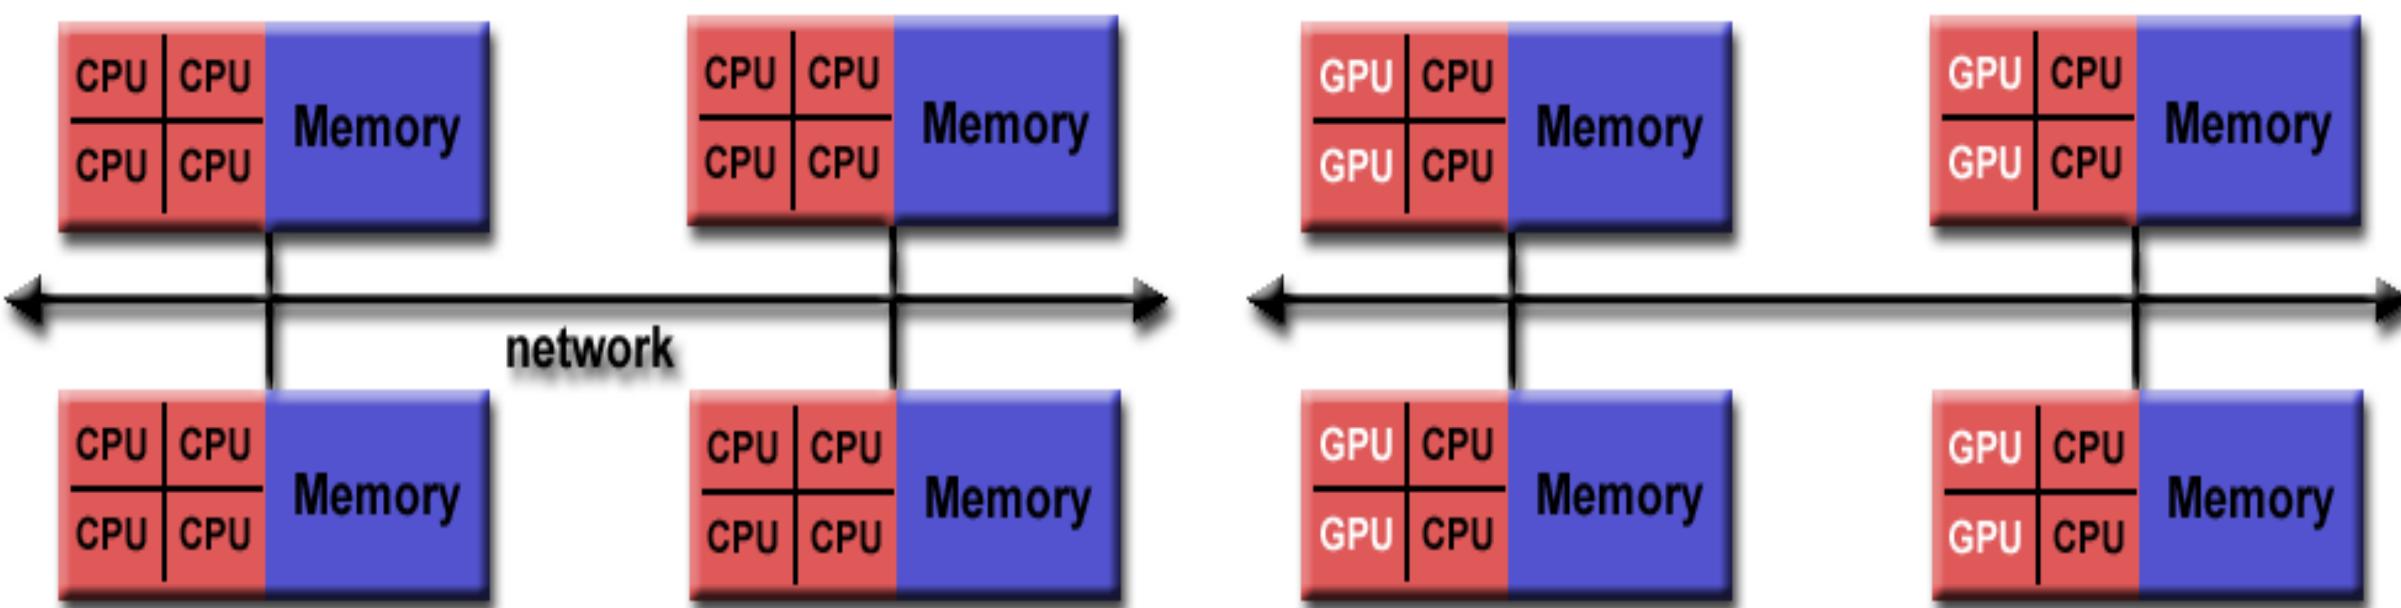


# Distributed Memory

- Advantages
  - Memory is scalable with the number of processors.
    - Increase the number of processors and the size of memory increases proportionately
  - Each processor can rapidly access its own memory without interference and without the overhead incurred with trying to maintain global cache coherency.
  - Cost effectiveness: can use commodity, off-the-shelf processors and networking
- Disadvantages
  - The programmer is responsible for many of the details associated with data communication between processors.
  - Non-uniform memory access times - data residing on a remote node takes longer to access than node local data.

# Hybrid Distributed-Shared Memory

- The largest and fastest computers in the world today employ both shared and distributed memory architectures
  - The shared memory component can be a shared memory machine or graphics processing units
  - The distributed memory component is the networking of multiple shared memory or GPU machines



# Hybrid Distributed-Shared Memory

- Advantages and Disadvantages
  - Whatever is common to both shared and distributed memory architectures
  - Increased scalability is an important advantage
  - Increased programmer complexity is an important disadvantage

# Parallel Programming Models

# Parallel Programming Model

- Programming model provides an **abstract view of computing system**
  - Abstraction above hardware and memory architectures
  - Value of a programming model is usually judged on its generality
    - how well a range of different problems can be expressed and
    - how well they execute on a range of different architectures
  - The implementation of a programming model can take several forms such as libraries invoked from traditional sequential languages, language extensions, or complete new execution models

# Parallel Programming Model

- Parallel programming models in common use:
  - Shared Memory (without threads)
  - Threads
  - Distributed Memory / Message Passing
  - Data Parallel
  - Hybrid
  - Single Program Multiple Data (SPMD)
  - Multiple Program Multiple Data (MPMD)
- These models are NOT specific to a particular type of machine or memory architecture
  - Any of these models can be implemented on any underlying hardware

# Parallel Programming Model

- SHARED memory model on a DISTRIBUTED memory machine
  - Machine memory was physically distributed across networked machines, but **appeared to the user as a single shared memory** (global address space).
  - This approach is referred to as **virtual shared memory**
- DISTRIBUTED memory model on a SHARED memory machine
  - The SGI Origin 2000 employed the CC-NUMA type of shared memory architecture, where every task has direct access to global address space spread across all machines.
  - However, the **ability to send and receive messages using MPI**, as is commonly done over a network of distributed memory machines, **was implemented and commonly used**

# Shared Memory Model - Without Threads

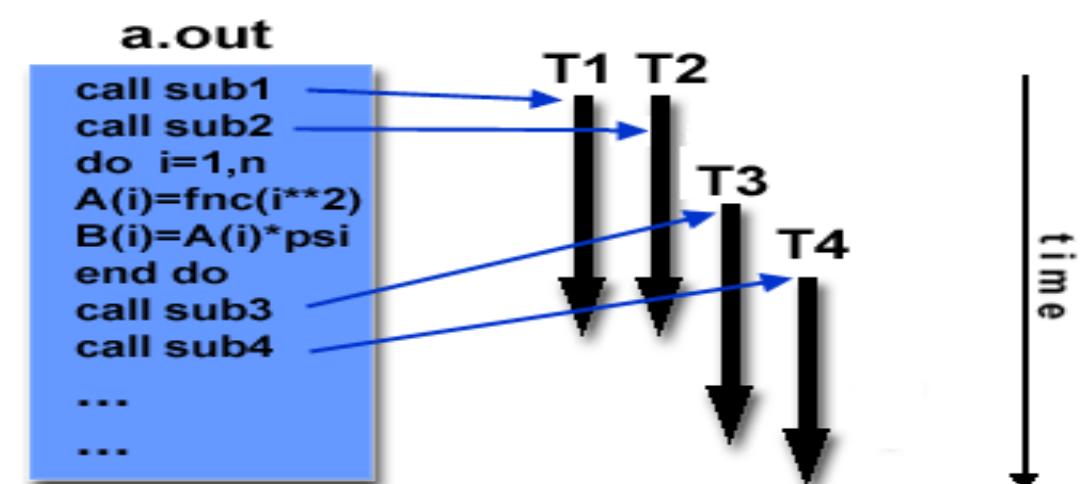
- Tasks share a common address space
  - Efficient means of passing data between programs
  - Various mechanisms such as locks / semaphores may be used to control access to the shared memory
  - Programmer's point of view
    - The notion of data "ownership" is lacking, so there is no **need to specify explicitly the communication of data between tasks**
    - Program development can often be simplified
- Disadvantage in terms of performance
  - It becomes more difficult to understand and manage data locality:
    - Keeping data local to the processor that works on it conserves memory accesses, cache refreshes and bus traffic that occurs when multiple processors use the same data

# Shared Memory Model - Without Threads

- Implementations
  - Native compilers or hardware translate user program variables into actual memory addresses, which are global
    - On stand-alone shared memory machines, this is straightforward.
    - On distributed shared memory machines, **memory is physically distributed across a network of machines**, but made global through specialized hardware and software

# Threads Model

- Type of shared memory programming model
  - A single "heavy weight" process can have multiple "light weight", concurrent execution paths
  - Main program a.out is scheduled by native OS
  - a.out loads and acquires all of the necessary system and user resources to run. This is the "heavy weight" process
  - a.out performs some serial work, and then creates a number of tasks (threads) that can be scheduled and run by the operating system concurrently



# Threads Model

- Each thread has local data, but also, **shares the entire resources of a.out**
- This saves the overhead associated with replicating a program's resources for each thread ("light weight"). Each thread also benefits from a global memory view because it shares the memory space of a.out
- **Threads communicate with each other through global memory** (updating address locations). This requires synchronization constructs to ensure that more than one thread is not updating the same global address at any time
- Threads can come and go, but a.out remains present to provide the necessary shared resources until the application has completed

# Threads Model

- Implementations

- *POSIX Threads*

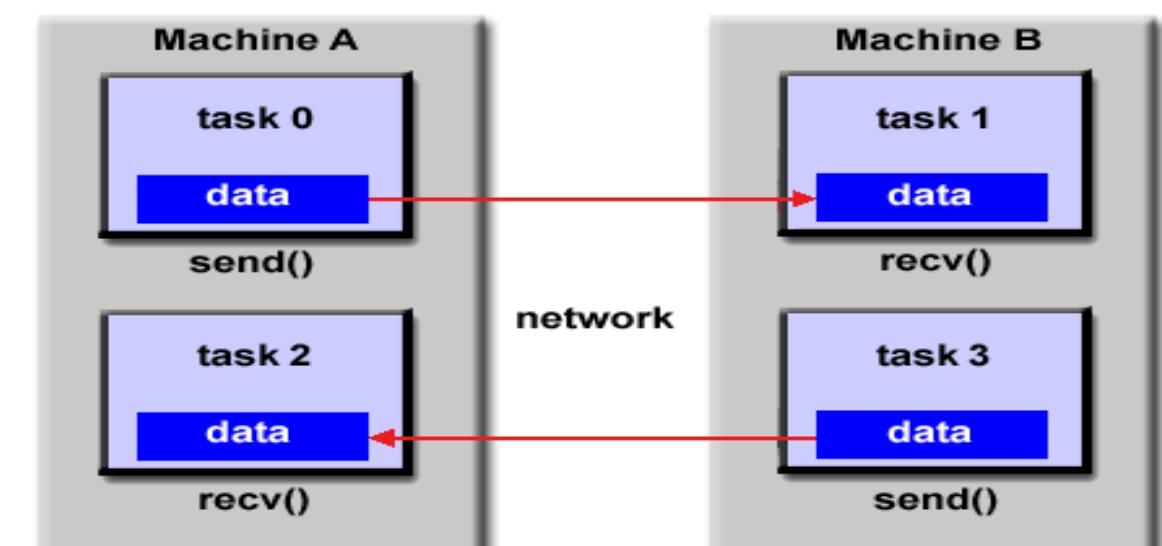
- Library based; requires parallel coding
    - C Language only
    - Commonly referred to as Pthreads.
    - Most hardware vendors now offer Pthreads in addition to their proprietary threads implementations.
    - Very explicit parallelism; requires significant programmer attention to detail.

- *OpenMP*

- Compiler directive based; can use serial code
    - Portable / multi-platform, including Unix and Windows platforms
    - Available in C/C++ and Fortran implementations
    - Can be very easy and simple to use

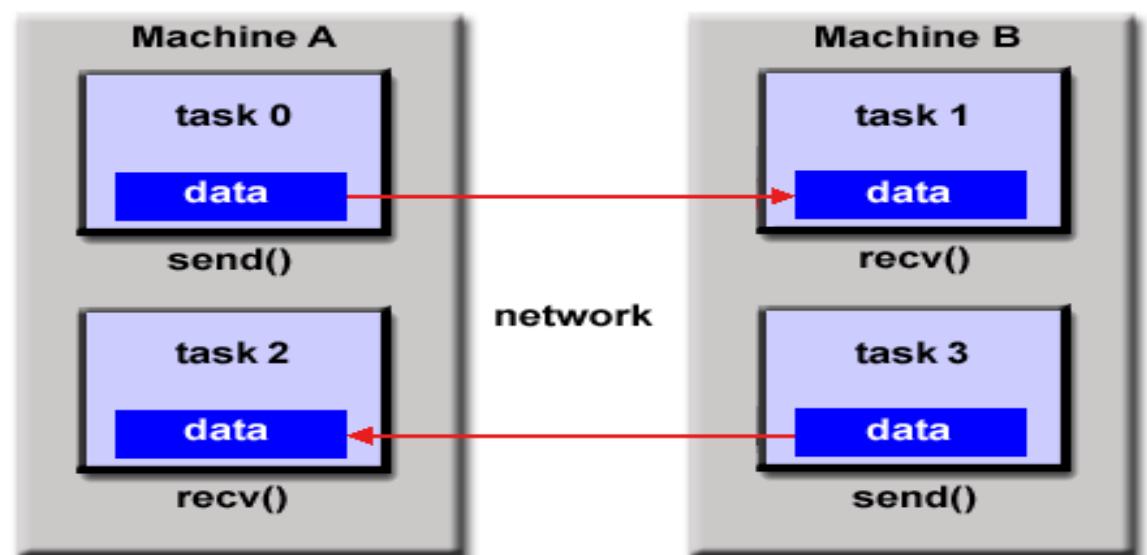
# Distributed Memory / Message Passing Model

- A set of tasks that use their own local memory during computation
  - Multiple tasks can reside on the same physical machine and/or across an arbitrary number of machines.
  - Tasks exchange data through communications by sending and receiving messages.
  - Data transfer usually requires cooperative operations to be performed by each process. For example, a send operation must have a matching receive operation.



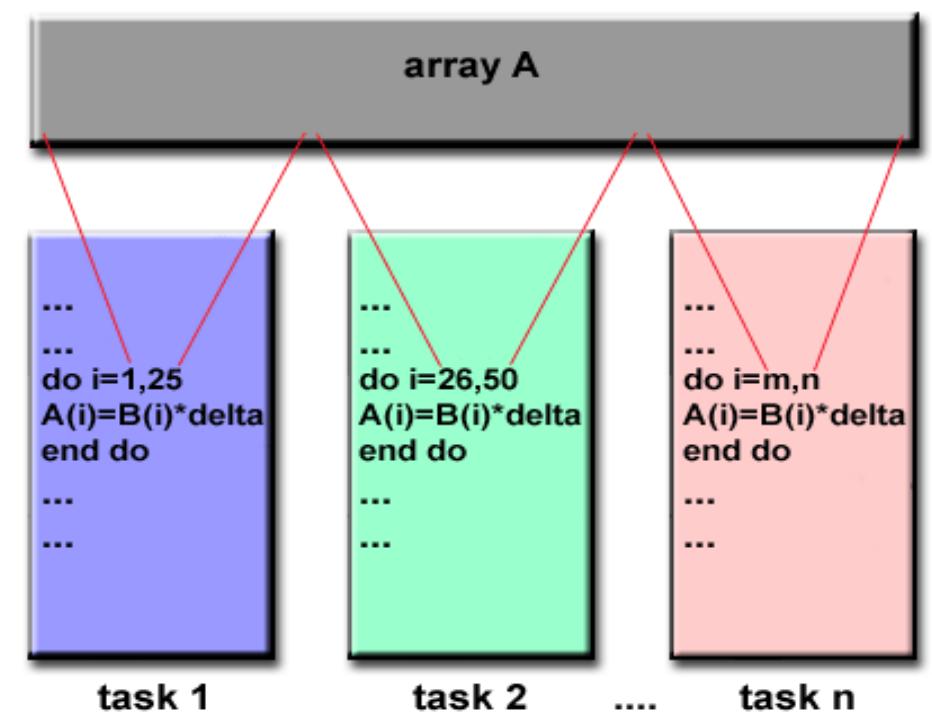
# Distributed Memory / Message Passing Model

- Implementations
  - From a programming perspective, message passing implementations usually comprise a library of subroutines
  - Calls to these subroutines are imbedded in source code
  - MPI specifications are available on the web at <http://www mpi-forum.org/docs/>
  - MPI implementations exist for virtually all popular parallel computing platforms



# Data Parallel Model

- Address space is treated globally
- A set of tasks work collectively on the same data structure, however, each task works on a different partition of the same data structure
  - On shared memory architectures, all tasks may have access to the data structure through global memory
  - On distributed memory architectures the data structure is split up and resides as "chunks" in the local memory of each task

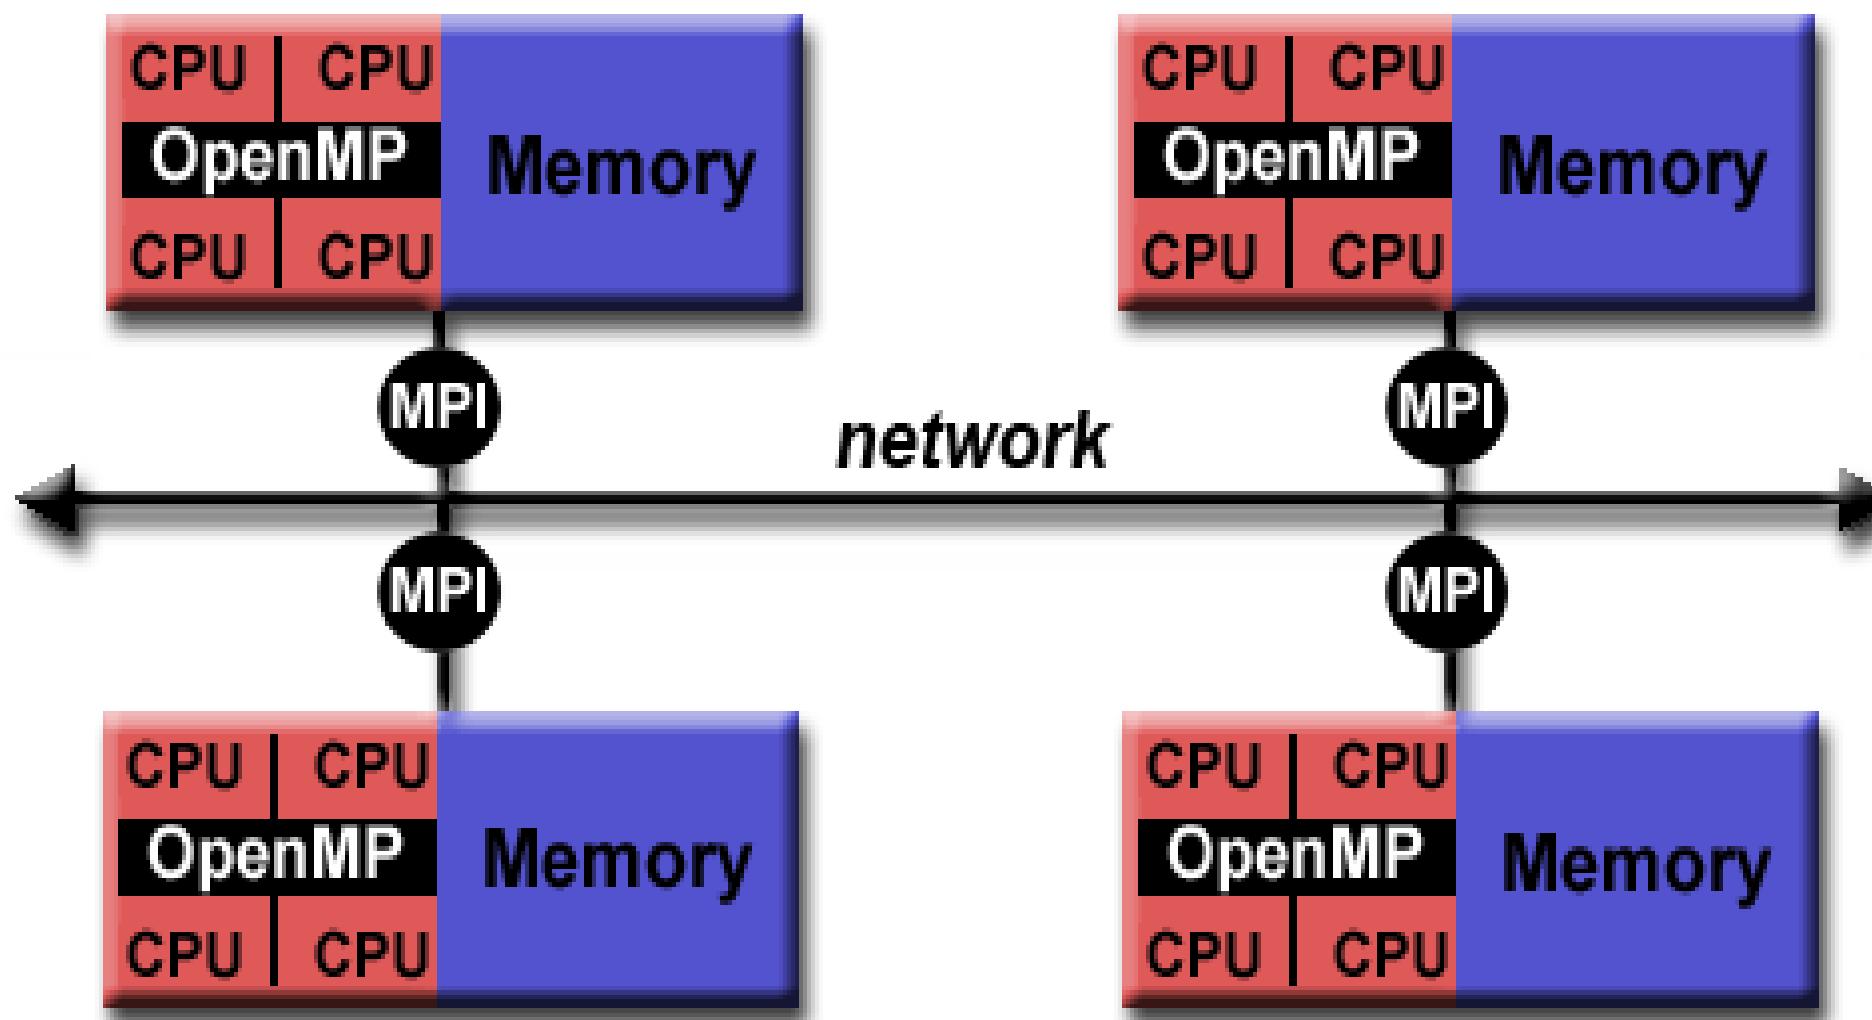


# Data Parallel Model

- Implementations
  - Unified Parallel C (UPC): an extension to the C programming language for SPMD parallel programming. Compiler dependent.  
More information: <http://upc.lbl.gov/>
  - Global Arrays: provides a shared memory style programming environment in the context of distributed array data structures. Public domain library with C and Fortran77 bindings. More information: <http://www.emsl.pnl.gov/docs/global/>
  - X10: a PGAS based parallel programming language being developed by IBM at the Thomas J. Watson Research Center. More information: <http://x10-lang.org/>

# Hybrid Model

- A hybrid model combines more than one of the previously described programming models



# Single Program Multiple Data

- High level programming model that can be built upon any combination of the previously mentioned parallel programming models
- **SINGLE PROGRAM:** All tasks execute their copy of the same program simultaneously.
  - This program can be threads, message passing, data parallel or hybrid.
- **MULTIPLE DATA:** All tasks may use different data



# Multiple Program Multiple Data

- High level programming model that can be built upon any combination of the previously mentioned parallel programming models
- MULTIPLE PROGRAM: Tasks may execute different programs simultaneously. The programs can be threads, message passing, data parallel or hybrid.
- MULTIPLE DATA: All tasks may use different data

